

Data on the Heap

num
bool
x char
x double

Next, lets add support for

- Data Structures

In the process of doing so, we will learn about

- Heap Allocation
- Run-time Tags
- High-order Func (Closures)

$\langle \text{env}, \text{code} \rangle$

Creating Heap Data Structures

We have already support for two primitive data types

```
data Ty
  = TNumber      -- e.g. 0,1,2,3,...
  | TBoolean     -- e.g. true, false
```

we could add several more of course, e.g.

- Char
- Double or Float

etc. (you should do it!)

However, for all of those, the same principle applies, more or less


- As long as the data fits into a single word (8-bytes)

Instead, lets learn how to make **unbounded data structures**

- Lists
- Trees
- ...

which require us to put data on the **heap**

not just the stack that we've used so far.





Stack vs. Heap

Pairs

While our *goal* is to get to lists and trees, the journey of a thousand miles begins with

a single step...

So! we will *begin* with the humble **pair**.

Construct

(2, 3) length=2

(1, (2, 3))

$e[0]$

$e[1]$

Assignment

$e[0]$

(1, 2, 3, 7, (9, 12))

"length" = ?
Storage

$e[i]$

Pairs: Semantics (Behavior)

First, let's ponder what exactly we're trying to achieve.

We want to enrich our language with *two* new constructs:

- **Constructing** pairs, with a new expression of the form (e_0, e_1) where e_0 and e_1 are expressions.
- **Accessing** pairs, with new expressions of the form $e[0]$ and $e[1]$ which

$e[0]$ $e[1]$

evaluate to the first and second element of the tuple e respectively.

For example,

```
let t = (2, 3) in  
  t[0] + t[1]
```

should evaluate to 5.

Strategy

Next, let's informally develop a strategy for extending our language with pairs, implementing the above semantics. We need to work out strategies for:

1. **Representing** pairs in the machine's memory,

$(e_0, e_1) \longrightarrow \langle asm \rangle$

$e[0]$ \longrightarrow $\langle \text{asm} \rangle$

2. **Constructing** pairs (i.e. implementing $(e0, e1)$ in assembly),
3. **Accessing** pairs (i.e. implementing $e[0]$ and $e[1]$ in assembly).

1. Representation

Recall that we represent all values: (05-cobra.md/#option-2-use-a-tag-bit)

- Number like 0, 1, 2 ...
- Boolean like true, false

} 64

as a ~~single~~ word either

- 8 bytes on the stack, or
- a single register rax, rbx etc.


EXERCISE

What kinds of problems do you think might arise if we represent a pair $(2, 3)$ on the *stack* as:

```
|   |
|---|
| 3 |
|---|
| 2 |
|---|
| ... |
|---|
```

Let $t = (\underline{2}, \underline{3}), 4)$ in

$t = (\underline{2}, 3), 4)$
 $t = (2, \underline{3}, 4)$



$t[0][0]$

1, 2, 3, 4, 5

$\text{cons}(1, \text{cons}(2, \text{cons}(3, \text{cons}(4, \text{nil}))))$
 $(1, (2, (3, (4, \text{false}))))$

$(1, 2) \quad 3$
 $\downarrow \quad \downarrow$
 (e_0, e_1)
 $e[0], e[1]$

QUIZ

How many words would we need to store the tuple

(3, (4, 5))

1. 1 word


2. 2 words

3. 3 words

4. 4 words

5. 5 words

l



tup

def tail(l):
 l[1]

def isNil():
 l == False

def nil():
 false

def cons(h, t):
 (h, t)

def range(lo, hi):
 if lo < hi:

 cons(lo, range(lo+1, hi))
 else:
 nil()

def length(l):
 if isNil(l):

 0
 else:
 1 + length(tail(l))


l = range(0, 100)
length(l)




Pointers

Every problem in computing can be solved by adding a level of indirection.

We will **represent a pair** by a **pointer** to a block of **two adjacent words** of memory.





Pairs on the heap

The above shows how the pair $(2, (3, (4, 5)))$ and its sub-pairs can be stored in the **heap** using pointers.

(4, 5) is stored by adjacent words storing

- 4 and
- 5

(3, (4, 5)) is stored by adjacent words storing

- 3 and
- a **pointer** to a heap location storing (4, 5)

(2, (3, (4, 5))) is stored by adjacent words storing

- 2 and
- a **pointer** to a heap location storing (3, (4, 5)).

A Problem: Numbers vs. Pointers?

How will we tell the difference between *numbers* and *pointers*?

That is, how can we tell the difference between

1. the *number* 5 and
2. a *pointer* to a block of memory (with address 5)?

Each of the above corresponds to a *different* tuple

1. (4, 5) or
2. (4, (...)).

so its pretty crucial that we have a way of knowing *which* value it is.

$t = (1, (2, 3))$


$t = 3$

$t = \text{false}$

Tagging Pointers

As you might have guessed, we can extend our tagging mechanism (05-cobra.md/#option-2-use-a-tag-bit) to account for *pointers*.

Type	LSB
number	xx0
boolean	111
pointer	001



$$t = (1, (2, 3))$$

That is, for

- number the **last bit** will be 0 (as before),
- boolean the **last 3 bits** will be 111 (as before), and
- pointer the **last 3 bits** will be 001.

Pointers are
8-byte aligned

(We have 3-bits worth for tags, so have wiggle room for other primitive types.)

Address Alignment

As we have a **3 bit tag**

- leaving **$64 - 3 = 61$ bits** for the actual address

So actual addresses, written in binary, omitting trailing zeros, are of the form


Binary	Decimal
0b000000000	0
0b00001000	8
0b00010000	16
0b00011000	24
0b00100000	32
...	

That is, the addresses are **8-byte aligned**.

Which is great because at each address, we have a pair, i.e. a **2-word = 16-byte block**, so the **next allocated address will also fall on an 8-byte boundary**.



- But ... what if we had 3-tuples? or 5-tuples? ...



tuple = pointer to heap with
 - tag set to 1
 →

2. Construction

Next, let's look at how to implement pair **construction** that is, generate the assembly for expressions like:

(e_1, e_2)

To construct a pair (e_1, e_2) we

1. **Allocate** a new 2-word block, and getting the starting address at rax ,
2. **Copy** the value of e_1 (resp. e_2) into $[rax]$ (resp. $[rax + 8]$).

3. **Tag** the last bit of `rax` with 1.

The resulting `eax` is the **value of the pair**

- The *last step* ensures that the value carries the proper tag.

ANF will ensure that `e1` and `e2` are immediate expressions (`04-b0a.md/#idea-immediate-expressions`)

- will make the second step above straightforward.

EXERCISE How will we do ANF conversion for $(e1, e2)$?

Allocating Addresses

Lets use a **global** register `r15` to maintain the address of the **next free block** on the heap.

Every time we need a *new* block, we will:

1. Copy the current `r15` into `rax`

- Set the last bit to 1 to ensure proper tagging.
- `rax` will be used to fill in the values

2. Increment the value of `r15` by 16

- Thus *allocating* 8 bytes (= 2 words) at the address in `rax`

Note that addresses stay 8-byte aligned (last 3 bits = 0) if we

- *Start* our blocks at an 8-byte boundary, and
- *Allocate* 16 bytes at a time,

NOTE: Your assignment will have *blocks of varying sizes*

- You will have to *maintain* the 8-byte alignment by *padding*


Example: Allocation

In the figure below, we have

- a source program on the left,
- the ANF equivalent next to it.




Source	ANF
<pre> let p = (3, (4, 5)) , x = p[0] = 3 , y = p[1][0] = 4 , z = p[1][1] = 5 in </pre>	<pre> 0 let anf0 = (4, 5) 1 , p = (3, anf0) 2 , x = p[0] , anf1 = p[1] y = anf1[0] </pre>



Example of Pairs

The figure below shows the how the heap and `r15` evolve at points 1, 2 and 3:




Allocating Pairs on the Heap

QUIZ

In the ANF version, `p` is the *second (local) variable* stored in the stack frame. What *value* gets moved into the *second stack slot* when evaluating the above program?

1. 0x3 \rightarrow 3
2. (3, (4, 5)) \rightarrow
3. 0x11 \rightarrow 17
4. 0x9 \rightarrow 9
5. 0x10 \rightarrow 16



3. Accessing

Finally, to **access** the elements of a pair

Lets compile $e[0]$ to get the first or $e[1]$ to get the second element


1. **Check** that immediate value e is a pointer
2. **Load** e into rbx
3. **Remove** the tag bit from rbx

4. Copy the value in `[rbx]` (resp. `[rbx + 8]`) into `rbx`.

Example: Access

Here is a snapshot of the heap after the pair(s) are allocated.

Source	ANF	Heap																		
<pre>let p = (3, (4, 5)) , x = p[0] , y = p[1][0] , z = p[1][1] in x + y + z</pre>	<pre>let anf0 = (4, 5) , p = (3, anf0) , x = p[0] , anf1 = p[1] , y = anf1[0] , z = anf1[1]</pre>	<table border="1"> <tr><td>0</td><td>0x8</td><td>4</td></tr> <tr><td>8</td><td>0xA</td><td>5</td></tr> <tr><td>16</td><td>0x6</td><td>3</td></tr> <tr><td>24</td><td>0x1</td><td></td></tr> <tr><td>32</td><td></td><td></td></tr> <tr><td>40</td><td></td><td></td></tr> </table>	0	0x8	4	8	0xA	5	16	0x6	3	24	0x1		32			40		
0	0x8	4																		
8	0xA	5																		
16	0x6	3																		
24	0x1																			
32																				
40																				



Allocating Pairs on the Heap

Lets work out how the values corresponding to x , y and z in the example above get stored on the stack frame in the course of evaluation.

Variable	Hex Value	Value
anf0	0x001	ptr 0
p	0x011	ptr 16
x	0x006	num 3
anf1	0x001	ptr 0
y	0x008	num 4
z	0x00A	num 5
anf2	0x00E	num 7
result	0x018	num 12